

# Lecture 2: Exploration and Exploitation

Hado van Hasselt

Reinforcement learning, 2021



# Background

Recommended reading:

Sutton & Barto 2018, Chapter 2


Further background material:

*Bandit Algorithms*, Lattimore & Szepesvári, 2020

*Finite-time analysis of the multiarmed bandit problem*, Auer, Cesa-Bianchi, Fischer, 2002



## Recap



- ▶ Reinforcement learning is the science of learning to make decisions
- ▶ Agents can learn a **policy**, **value function** and/or a **model**
- ▶ The general problem involves taking into account **time** and **consequences**
- ▶ Decisions affect the **reward**, the **agent state**, and **environment state**
- ▶ Learning is **active**: decisions impact data



# This Lecture

In this lecture, we simplify the setting

- ▶ The environment is assumed to have only a **single state**
- ▶  $\implies$  actions no longer have long-term consequences in the environment
- ▶  $\implies$  actions still do impact **immediate reward**
- ▶  $\implies$  other observations can be ignored
- ▶ We discuss how to learn a policy in this setting



# Blackboard: Example



# Exploration vs. Exploitation

- ▶ Learning agents need to trade off two things
  - ▶ **Exploitation**: Maximise performance based on current knowledge
  - ▶ **Exploration**: Increase knowledge
- ▶ We need to gather information to make the best overall decisions
- ▶ The best long-term strategy may involve short-term sacrifices



# Formalising the problem



# The Multi-Armed Bandit

- ▶ A multi-armed bandit is a set of distributions  $\{\mathcal{R}_a | a \in \mathcal{A}\}$
- ▶  $\mathcal{A}$  is a (known) set of actions (or “arms”)
- ▶  $\mathcal{R}_a$  is a distribution on rewards, given action  $a$
- ▶ At each step  $t$  the agent selects an action  $A_t \in \mathcal{A}$
- ▶ The environment generates a reward  $R_t \sim \mathcal{R}_{A_t}$
- ▶ The goal is to maximise cumulative reward  $\sum_{i=1}^t R_i$
- ▶ We do this by learning a **policy**: a distribution on  $\mathcal{A}$





# Values and Regret

- ▶ The **action value** for action  $a$  is the expected reward

$$q(a) = \mathbb{E} [R_t | A_t = a]$$

- ▶ The **optimal value** is

$$v_* = \max_{a \in \mathcal{A}} q(a) = \max_a \mathbb{E} [R_t | A_t = a]$$

- ▶ **Regret** of an action  $a$  is

$$\Delta_a = v_* - q(a)$$

- ▶ The regret for the optimal action is zero



# Regret

- ▶ We want to minimise **total regret**:

$$L_t = \sum_{n=1}^t v_* - q(A_n) = \sum_{n=1}^t \Delta_{A_n}$$

- ▶ Maximise cumulative reward  $\equiv$  minimise total regret
- ▶ The summation spans over the full ‘lifetime of learning’



# Algorithms



# Algorithms

- ▶ We will discuss several algorithms:
  - ▶ Greedy
  - ▶  $\epsilon$ -greedy
  - ▶ UCB
  - ▶ Thompson sampling
  - ▶ Policy gradients
- ▶ The first three all use **action value estimates**  $Q_t(a) \approx q(a)$



# Action values

- ▶ The **action value** for action  $a$  is the expected reward

$$q(a) = \mathbb{E}[R_t | A_t = a]$$

- ▶ A simple estimate is the average of the sampled rewards:

$$Q_t(a) = \frac{\sum_{n=1}^t \mathcal{I}(A_n = a) R_n}{\sum_{n=1}^t \mathcal{I}(A_n = a)}$$

$\mathcal{I}(\cdot)$  is the **indicator** function:  $\mathcal{I}(\text{True}) = 1$  and  $\mathcal{I}(\text{False}) = 0$

- ▶ The **count** for action  $a$  is

$$N_t(a) = \sum_{n=1}^t \mathcal{I}(A_n = a)$$



## Action values

- ▶ This can also be updated incrementally:

$$Q_t(A_t) = Q_{t-1}(A_t) + \alpha_t \underbrace{(R_t - Q_{t-1}(A_t))}_{\text{error}},$$

$$\forall a \neq A_t : Q_t(a) = Q_{t-1}(a)$$

with

$$\alpha_t = \frac{1}{N_t(A_t)} \quad \text{and} \quad N_t(A_t) = N_{t-1}(A_t) + 1,$$

where  $N_0(a) = 0$ .

- ▶ We will later consider other **step sizes**  $\alpha$
- ▶ For instance, constant  $\alpha$  would lead to **tracking**, rather than averaging



Algorithms: greedy



# The greedy policy

- ▶ One of the simplest policies is **greedy**:
  - ▶ Select action with highest value:  $A_t = \operatorname{argmax}_a Q_t(a)$
  - ▶ Equivalently:  $\pi_t(a) = \mathcal{I}(A_t = \operatorname{argmax}_a Q_t(a))$  (assuming no ties are possible)





# Example:

## Regret of the greedy policy





Algorithms:  $\epsilon$ -greedy



# $\epsilon$ -Greedy Algorithm

- ▶ Greedy can get stuck on a suboptimal action forever  
 $\implies$  linear expected total regret
- ▶ The  $\epsilon$ -greedy algorithm:
  - ▶ With probability  $1 - \epsilon$  select greedy action:  $a = \operatorname{argmax}_{a \in \mathcal{A}} Q_t(a)$
  - ▶ With probability  $\epsilon$  select a random action
  - ▶ Equivalently:

$$\pi_t(a) = \begin{cases} (1 - \epsilon) + \epsilon/|\mathcal{A}| & \text{if } Q_t(a) = \max_b Q_t(b) \\ \epsilon/|\mathcal{A}| & \text{otherwise} \end{cases}$$

- ▶  $\epsilon$ -greedy continues to explore  
 $\implies \epsilon$ -greedy with constant  $\epsilon$  has linear expected total regret



# Algorithms: Policy gradients



# Policy search

- ▶ Can we learn policies  $\pi(a)$  directly, instead of learning values?
- ▶ For instance, define **action preferences**  $H_t(a)$  and a policy

$$\pi(a) = \frac{e^{H_t(a)}}{\sum_b e^{H_t(b)}} \quad (\text{softmax})$$

- ▶ The preferences are not values: they are just learnable policy parameters
- ▶ Goal: learn by optimising the preferences



# Policy gradients

- ▶ Idea: update policy parameters such that expected value increases
- ▶ We can use **gradient ascent**
- ▶ In the bandit case, we want to update:

$$\theta_{t+1} = \theta_t + \alpha \nabla_{\theta} \mathbb{E}[R_t | \pi_{\theta_t}],$$

where  $\theta_t$  are the current policy parameters

- ▶ Can we compute this gradient?



## Gradient bandits

- Log-likelihood trick (also known as REINFORCE trick, Williams 1992):

$$\begin{aligned}\nabla_{\theta} \mathbb{E}[R_t | \pi_{\theta}] &= \nabla_{\theta} \sum_a \pi_{\theta}(a) \overbrace{\mathbb{E}[R_t | A_t = a]}^{= q(a)} \\&= \sum_a q(a) \nabla_{\theta} \pi_{\theta}(a) \\&= \sum_a q(a) \frac{\pi_{\theta}(a)}{\pi_{\theta}(a)} \nabla_{\theta} \pi_{\theta}(a) \\&= \sum_a \pi_{\theta}(a) q(a) \frac{\nabla_{\theta} \pi_{\theta}(a)}{\pi_{\theta}(a)} \\&= \mathbb{E} \left[ R_t \frac{\nabla_{\theta} \pi_{\theta}(A_t)}{\pi_{\theta}(A_t)} \right] \qquad = \mathbb{E} [R_t \nabla_{\theta} \log \pi_{\theta}(A_t)]\end{aligned}$$





# Gradient bandits

- ▶ Log-likelihood trick (also known as REINFORCE trick, Williams 1992):

$$\nabla_{\theta} \mathbb{E}[R_t | \theta] = \mathbb{E}[R_t \nabla_{\theta} \log \pi_{\theta}(A_t)]$$

- ▶ We can sample this!
- ▶ So

$$\theta = \theta + \alpha R_t \nabla_{\theta} \log \pi_{\theta}(A_t),$$

this is **stochastic gradient ascent** on the (true) value of the policy

- ▶ Can use **sampled** rewards — does not need value estimates



# Gradient bandits

- For soft max:

$$\begin{aligned}H_{t+1}(a) &= H_t(a) + \alpha R_t \frac{\partial \log \pi_t(A_t)}{\partial H_t(a)} \\&= H_t(a) + \alpha R_t (\mathcal{I}(a = A_t) - \pi_t(a))\end{aligned}$$

- $\Rightarrow$

$$\begin{aligned}H_{t+1}(A_t) &= H_t(A_t) + \alpha R_t (1 - \pi_t(A_t)) \\H_{t+1}(a) &= H_t(a) - \alpha R_t \pi_t(a) \quad \text{if } a \neq A_t\end{aligned}$$

- Preferences for actions with higher rewards increase more (or decrease less), making them more likely to be selected again



Theory: what is possible?



# How well can we do?

## Theorem (Lai and Robbins)

*Asymptotic total regret is at least logarithmic in number of steps*

$$\lim_{t \rightarrow \infty} L_t \geq \log t \sum_{a | \Delta_a > 0} \frac{\Delta_a}{KL(\mathcal{R}_a || \mathcal{R}_{a^*})}$$

(Note:  $KL(\mathcal{R}_a || \mathcal{R}_{a^*}) \propto \Delta_a^2$ )

- ▶ Note that **regret grows at least logarithmically**
- ▶ That's still a whole lot better than linear growth! Can we get it in practice?
- ▶ Are there algorithms for which the **upper bound** is logarithmic as well?



# Counting Regret

- ▶ Recall  $\Delta_a = v_* - q(a)$
- ▶ Total regret depends on action regrets  $\Delta_a$  and action counts

$$L_t = \sum_{n=1}^t \Delta_{A_n} = \sum_{a \in \mathcal{A}} N_t(a) \Delta_a$$


- ▶ A good algorithm ensures small counts for large action regrets



Optimism in the face of uncertainty




# Optimism in the Face of Uncertainty



- ▶ Which action should we pick?
- ▶ More uncertainty about its value: more important to explore that action




## Optimism in the Face of Uncertainty







## Optimism in the Face of Uncertainty




## Optimism in the Face of Uncertainty




## Optimism in the Face of Uncertainty




## Optimism in the Face of Uncertainty




## Optimism in the Face of Uncertainty



## Optimism in the Face of Uncertainty



## Optimism in the Face of Uncertainty



# Algorithms: UCB





# Upper Confidence Bounds

- ▶ Estimate an upper confidence  $U_t(a)$  for each action value, such that  $q(a) \leq Q_t(a) + U_t(a)$  with high probability
- ▶ Select action maximizing **upper confidence bound** (UCB)

$$a_t = \operatorname{argmax}_{a \in \mathcal{A}} Q_t(a) + U_t(a)$$

- ▶ The uncertainty should depend on the number of times  $N_t(a)$  action  $a$  has been selected
  - ▶ Small  $N_t(a) \Rightarrow$  large  $U_t(a)$  (estimated value is uncertain)
  - ▶ Large  $N_t(a) \Rightarrow$  small  $U_t(a)$  (estimated value is accurate)
- ▶ Then  $a$  is only selected if either...
  - ▶ ... $Q_t(a)$  is large (=good action), or
  - ▶ ... $U_t(a)$  is large (=high uncertainty) (or both)
- ▶ Can we **derive** an optimal bound?



# Theory: the optimality of UCB



# Hoeffding's Inequality

## Theorem (Hoeffding's Inequality)

Let  $X_1, \dots, X_n$  be i.i.d. random variables in  $[0, 1]$  with true mean  $\mu = \mathbb{E}[X]$ , and let  $\bar{X}_n = \frac{1}{n} \sum_{i=1}^n X_i$  be the sample mean. Then

$$p\left(\bar{X}_n + u \leq \mu\right) \leq e^{-2nu^2}$$

- ▶ We can apply Hoeffding's Inequality to bandits with bounded rewards
- ▶ If  $R_t \in [0, 1]$ , then

$$p(Q_t(a) + U_t(a) \leq q(a)) \leq e^{-2N_t(a)U_t(a)^2}$$

- ▶ By symmetry, we can also flip it around

$$p(Q_t(a) - U_t(a) \geq q(a)) \leq e^{-2N_t(a)U_t(a)^2}$$



# Calculating Upper Confidence Bounds

- ▶ We can pick a maximal desired probability  $p$  that the true value exceeds an upper bound and solve for this bound  $U_t(a)$

$$e^{-2N_t(a)U_t(a)^2} = p$$
$$\implies U_t(a) = \sqrt{\frac{-\log p}{2N_t(a)}}$$

We then know the probability that this happens is smaller than  $p$

- ▶ Idea: reduce  $p$  as we observe more rewards, e.g.,  $p = 1/t$

$$U_t(a) = \sqrt{\frac{\log t}{2N_t(a)}}$$

- ▶ This ensures that we always keep exploring, but not too much



- ▶ UCB:

$$a_t = \operatorname{argmax}_{a \in \mathcal{A}} Q_t(a) + c \sqrt{\frac{\log t}{N_t(a)}}$$

- ▶ Intuition:
  - ▶ If  $\Delta_a$  is large, then  $N_t(a)$  is small, because  $Q_t(a)$  is likely to be small
  - ▶ So either  $\Delta_a$  is small or  $N_t(a)$  is small
  - ▶ In fact, we can prove  $\Delta_a N_t(a) \leq O(\log t)$ , for all  $a$



# UCB

► UCB:

$$a_t = \operatorname{argmax}_{a \in \mathcal{A}} Q_t(a) + c \sqrt{\frac{\log t}{N_t(a)}}$$

where  $c$  is a hyper-parameter

**Theorem (Auer et al., 2002)**

*UCB with  $c = \sqrt{2}$  achieves logarithmic expected total regret*

$$L_t \leq 8 \sum_{a | \Delta_a > 0} \frac{\log t}{\Delta_a} + O\left(\sum_a \Delta_a\right), \quad \forall t.$$



# Blackboard: UCB derivation







# Bayesian approaches



# Bayesian Bandits

- ▶ We could adopt **Bayesian** approach and model distributions over values  $p(q(a) \mid \theta_t)$
- ▶ This is interpreted as our **belief** that, e.g.,  $q(a) = x$  for all  $x \in \mathbb{R}$
- ▶ E.g.,  $\theta_t$  could contain the means and variances of Gaussian belief distributions
- ▶ Allows us to inject rich prior knowledge  $\theta_0$
- ▶ We can then use posterior belief to guide exploration




# Bayesian Bandits: Example

- ▶ Consider bandits with **Bernoulli** reward distribution: rewards are 0 or +1
- ▶ For each action, the prior could be a **uniform distribution** on  $[0, 1]$
- ▶ This means we think each value in  $[0, 1]$  is equally likely
- ▶ The posterior is a Beta distribution  $\text{Beta}(x_a, y_a)$  with initial parameters  $x_a = 1$  and  $y_a = 1$  for each action  $a$
- ▶ Updating the posterior:
  - ▶  $x_{A_t} \leftarrow x_{A_t} + 1$  when  $R_t = 0$
  - ▶  $y_{A_t} \leftarrow y_{A_t} + 1$  when  $R_t = 1$




## Bayesian Bandits: Example

Suppose:  $R_1 = +1, R_2 = +1, R_3 = 0, R_4 = 0$



# Bayesian Bandits with Upper Confidence Bounds



- ▶ We can estimate upper confidences from the posterior
  - ▶ e.g.,  $U_t(a) = c\sigma_t(a)$  where  $\sigma(a)$  is std dev of  $p_t(q(a))$
- ▶ Then, pick an action that maximises  $Q_t(a) + c\sigma(a)$



# Algorithms: Thompson sampling



# Probability Matching

- ▶ A different option is to use **probability matching**:  
Select action  $a$  according to the probability (belief) that  $a$  is optimal

$$\pi_t(a) = p\left(q(a) = \max_{a'} q(a') \mid \mathcal{H}_{t-1}\right)$$

- ▶ Probability matching is optimistic in the face of uncertainty:  
Actions have higher probability when either the estimated value is high, or the uncertainty is high
- ▶ Can be difficult to compute  $\pi(a)$  analytically from posterior (but can be done numerically)



# Thompson Sampling

- ▶ Thompson sampling (Thompson 1933):
  - ▶ Sample  $Q_t(a) \sim p_t(q(a))$ ,  $\forall a$
  - ▶ Select action maximising sample,  $A_t = \operatorname{argmax}_{a \in \mathcal{A}} Q_t(a)$
- ▶ **Thompson sampling** is sample-based probability matching

$$\begin{aligned}\pi_t(a) &= \mathbb{E} \left[ \mathcal{I}(Q_t(a) = \max_{a'} Q_t(a')) \right] \\ &= p \left( q(a) = \max_{a'} q(a') \right)\end{aligned}$$

- ▶ For Bernoulli bandits, Thompson sampling achieves Lai and Robbins lower bound on regret, and therefore is **optimal**





Planning to explore



# Information State Space

- ▶ We have viewed bandits as **one-step** decision-making problems
- ▶ Can also view as **sequential** decision-making problems
- ▶ Each step the agent updates state  $S_t$  to summarise the past
- ▶ Each action  $A_t$  causes a transition to a new **information state**  $S_{t+1}$  (by adding information), with probability  $p(S_{t+1} \mid A_t, S_t)$
- ▶ We now have a Markov decision problem
- ▶ The state is fully internal to the agent
- ▶ State transitions are random due to rewards & actions
- ▶ Even in bandits actions affect the future after all, via learning



## Example: Bernoulli Bandits

- ▶ Consider a Bernoulli bandit, such that

$$p(R_t = 1 \mid A_t = a) = \mu_a$$

$$p(R_t = 0 \mid A_t = a) = 1 - \mu_a$$

- ▶ E.g., win or lose a game with probability  $\mu_a$
- ▶ Want to find which arm has the highest  $\mu_a$
- ▶ The information state is  $I = (\alpha, \beta)$ 
  - ▶  $\alpha_a$  counts the pulls of arm  $a$  where reward was 0
  - ▶  $\beta_a$  counts the pulls of arm  $a$  where reward was 1



# Solving Information State Space Bandits

- ▶ We formulated the bandit as an infinite MDP over information states
- ▶ This can be solved by reinforcement learning
- ▶ E.g., learn a Bayesian reward distribution, plan into the future
- ▶ This is known as **Bayes-adaptive** RL:  
optimally trades off exploration with respect to the prior distribution
- ▶ Can be extended to full RL, by also learning a transition model
- ▶ Can be unwieldy... unclear how to scale effectively



Example



End of lecture

